

2cs317-pes2ug22cs325-pes2ug22cs910

November 18, 2023

DATASET-2: Diamond Prices

Features Description: • Price: price in US dollars

- Carat: is the diamond's physical weight measured in metric carats.
- Cut: quality of the cut
- Color: diamond color, from J (worst) to D (best)
- Clarity: a measurement of how clear the diamond is
- X: length in mm
- Y: width in mm
- Z: depth in mm
- Depth: total depth percentage = $z / \text{mean}(x, y) = 2 * z / (x + y)$

Table: width of the top of diamond relative to widest poin

```
[24]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# Read the dataset
df = pd.read_csv('dataset_2_Diamonds Prices.csv')
print(df)
```

	index	carat	cut	color	clarity	depth	table	price	x	y	\
0	1	0.23	Ideal	E	SI2	61.5	55.0	326	3.95	3.98	
1	2	0.21	Premium	E	SI1	59.8	61.0	326	3.89	3.84	
2	3	0.23	Good	E	VS1	56.9	65.0	327	4.05	4.07	
3	4	0.29	Premium	I	VS2	62.4	58.0	334	4.20	4.23	
4	5	0.31	Good	J	SI2	63.3	58.0	335	4.34	4.35	

	z
0	2.43
1	2.31
2	2.31
3	2.63

```
4      2.75
...    ...
53938  3.74
53939  3.64
53940  3.49
53941  3.43
53942  3.47
```

```
[53943 rows x 11 columns]
```

Questions: 1.

- i) What is the shape of the dataset? (Specify rows and columns separately)
- ii) List the column names and their data types?
- iii) Delete 'index' column?

```
[25]: # Get the shape of the DataFrame
shape = df.shape
# Print the shape
print(f'The dataset has {shape[0]} rows and {shape[1]} columns.')
print()

# list of columns names and their data types
colTy = df.dtypes
# Print the column names and their data types
print(colTy)
print()

# remove the index column and print the new DataFrame
coldel = 'index'
df = df.drop(columns=[coldel])
print(df)
```

The dataset has 53943 rows and 11 columns.

```
index      int64
carat      float64
cut         object
color       object
clarity     object
depth      float64
table      float64
price      int64
x           float64
y           float64
z           float64
dtype: object
```

	carat	cut	color	clarity	depth	table	price	x	y	z
0	0.23	Ideal	E	SI2	61.5	55.0	326	3.95	3.98	2.43
1	0.21	Premium	E	SI1	59.8	61.0	326	3.89	3.84	2.31
2	0.23	Good	E	VS1	56.9	65.0	327	4.05	4.07	2.31
3	0.29	Premium	I	VS2	62.4	58.0	334	4.20	4.23	2.63
4	0.31	Good	J	SI2	63.3	58.0	335	4.34	4.35	2.75
...
53938	0.86	Premium	H	SI2	61.0	58.0	2757	6.15	6.12	3.74
53939	0.75	Ideal	D	SI2	62.2	55.0	2757	5.83	5.87	3.64
53940	0.71	Premium	E	SI1	60.5	55.0	2756	5.79	5.74	3.49
53941	0.71	Premium	F	SI1	59.8	62.0	2756	5.74	5.73	3.43
53942	0.70	Very Good	E	VS2	60.5	59.0	2757	5.71	5.76	3.47

[53943 rows x 10 columns]

Question 2.

Describe the summary statistics, min, max, mean, standard deviation for all numeric columns?

```
[26]: summary = df.describe()
print(summary)

print()
minimum = summary.min()
print("Min:\n",minimum)

print()
maximum = summary.max()
print("Max:\n",maximum)

print()
mean = summary.loc['mean']
print("Mean:\n",mean)

print()
standard_deviation = summary.loc['std']
print("Standard Deviation:\n",standard_deviation)
```

	carat	depth	table	price	x \
count	53943.000000	53935.000000	53943.000000	53943.000000	53931.000000
mean	0.797935	61.749426	57.457251	3932.734294	5.731166
std	0.473999	1.432672	2.234549	3989.338447	1.121819
min	0.200000	43.000000	43.000000	326.000000	0.000000
25%	0.400000	61.000000	56.000000	950.000000	4.710000
50%	0.700000	61.800000	57.000000	2401.000000	5.700000
75%	1.040000	62.500000	59.000000	5324.000000	6.540000
max	5.010000	79.000000	95.000000	18823.000000	10.740000

y

z

count	53934.000000	53933.000000
mean	5.734518	3.538768
std	1.142165	0.705728
min	0.000000	0.000000
25%	4.720000	2.910000
50%	5.710000	3.530000
75%	6.540000	4.040000
max	58.900000	31.800000

Min:

carat	0.200000
depth	1.432672
table	2.234549
price	326.000000
x	0.000000
y	0.000000
z	0.000000

dtype: float64

Max:

carat	53943.0
depth	53935.0
table	53943.0
price	53943.0
x	53931.0
y	53934.0
z	53933.0

dtype: float64

Mean:

carat	0.797935
depth	61.749426
table	57.457251
price	3932.734294
x	5.731166
y	5.734518
z	3.538768

Name: mean, dtype: float64

Standard Deviation:

carat	0.473999
depth	1.432672
table	2.234549
price	3989.338447
x	1.121819
y	1.142165
z	0.705728

Name: std, dtype: float64

Question 3.

List all distinct values and most frequent values in each column 'cut', 'colour' and 'clarity'?

```
[27]: # Assuming df is your DataFrame

# List all distinct values
cut_unique = df['cut'].unique()
color_unique = df['color'].unique()
clarity_unique = df['clarity'].unique()

# Get the most frequent values
cut_most_frequent = df['cut'].mode()[0]
color_most_frequent = df['color'].mode()[0]
clarity_most_frequent = df['clarity'].mode()[0]

# Print all distinct values
print(f"Distinct values in 'cut': {cut_unique}")
print(f"Distinct values in 'color': {color_unique}")
print(f"Distinct values in 'clarity': {clarity_unique}")

# Print the most frequent values
print(f"Most frequent value in 'cut': {cut_most_frequent}")
print(f"Most frequent value in 'color': {color_most_frequent}")
print(f"Most frequent value in 'clarity': {clarity_most_frequent}")
```

```
Distinct values in 'cut': ['Ideal' 'Premium' 'Good' 'Very Good' 'Fair']
Distinct values in 'color': ['E' 'I' 'J' 'H' 'F' 'G' 'D']
Distinct values in 'clarity': ['SI2' 'SI1' 'VS1' 'VS2' 'VVS2' 'VVS1' 'I1' 'IF']
Most frequent value in 'cut': Ideal
Most frequent value in 'color': G
Most frequent value in 'clarity': SI1
```

Question 4:

Identify and describe any data quality issues or inconsistencies within the data set. What steps would you take to clean and pre-processes the data to ensure its accuracy for further analysis

To clean and preprocess the data:

1. Handle missing values
2. Remove duplicates
3. Handle outliers
4. Convert data types

```
[28]: # Identify missing values
print("Missing values\n",df.isnull().sum())
# Identify duplicated records
print("Duplicated records\n",df.duplicated().sum())
# Identify outliers
```

```
print("Outliers\n",df.describe(include='all').loc['max'])
# Indetify inconsistent data types
print("Inconsistent data types\n",df.dtypes)
```

Missing values

```
carat      0
cut         0
color       0
clarity     0
depth       8
table       0
price       0
x           12
y           9
z           10
```

dtype: int64

Duplicated records

```
149
```

Outliers

```
carat      5.01
cut         NaN
color       NaN
clarity     NaN
depth      79.0
table      95.0
price     18823.0
x          10.74
y          58.9
z          31.8
```

Name: max, dtype: object

Inconsistent data types

```
carat      float64
cut         object
color       object
clarity     object
depth      float64
table      float64
price       int64
x          float64
y          float64
z          float64
```

dtype: object

Question 5:

- (i) Convert price in us dollar to rupees? (1 dollar = 80 rupees)
- (ii) Create a new column called 'color_clarity_cut' and values are color+ ' ' +clarity+ ' ' + cut? (Ex: E_ SI2_Ideal , E_ SI1_Premium)

```
[29]: # (i) Convert price in us dollar to rupees
df['price'] = df['price'] * 80
print(df)

# (ii) Create a new column called 'color_clarity_cut'
df['color_clarity_cut'] = df['color'] + '_' + df['clarity'] + '_' + df['cut']
print(df)
```

	carat	cut	color	clarity	depth	table	price	x	y	z
0	0.23	Ideal	E	SI2	61.5	55.0	26080	3.95	3.98	2.43
1	0.21	Premium	E	SI1	59.8	61.0	26080	3.89	3.84	2.31
2	0.23	Good	E	VS1	56.9	65.0	26160	4.05	4.07	2.31
3	0.29	Premium	I	VS2	62.4	58.0	26720	4.20	4.23	2.63
4	0.31	Good	J	SI2	63.3	58.0	26800	4.34	4.35	2.75
...
53938	0.86	Premium	H	SI2	61.0	58.0	220560	6.15	6.12	3.74
53939	0.75	Ideal	D	SI2	62.2	55.0	220560	5.83	5.87	3.64
53940	0.71	Premium	E	SI1	60.5	55.0	220480	5.79	5.74	3.49
53941	0.71	Premium	F	SI1	59.8	62.0	220480	5.74	5.73	3.43
53942	0.70	Very Good	E	VS2	60.5	59.0	220560	5.71	5.76	3.47

[53943 rows x 10 columns]

	carat	cut	color	clarity	depth	table	price	x	y	z	\
0	0.23	Ideal	E	SI2	61.5	55.0	26080	3.95	3.98	2.43	
1	0.21	Premium	E	SI1	59.8	61.0	26080	3.89	3.84	2.31	
2	0.23	Good	E	VS1	56.9	65.0	26160	4.05	4.07	2.31	
3	0.29	Premium	I	VS2	62.4	58.0	26720	4.20	4.23	2.63	
4	0.31	Good	J	SI2	63.3	58.0	26800	4.34	4.35	2.75	
...	
53938	0.86	Premium	H	SI2	61.0	58.0	220560	6.15	6.12	3.74	
53939	0.75	Ideal	D	SI2	62.2	55.0	220560	5.83	5.87	3.64	
53940	0.71	Premium	E	SI1	60.5	55.0	220480	5.79	5.74	3.49	
53941	0.71	Premium	F	SI1	59.8	62.0	220480	5.74	5.73	3.43	
53942	0.70	Very Good	E	VS2	60.5	59.0	220560	5.71	5.76	3.47	

	color_clarity_cut
0	E_SI2_Ideal
1	E_SI1_Premium
2	E_VS1_Good
3	I_VS2_Premium
4	J_SI2_Good
...	...
53938	H_SI2_Premium
53939	D_SI2_Ideal
53940	E_SI1_Premium
53941	F_SI1_Premium
53942	E_VS2_Very Good

[53943 rows x 11 columns]

Question 6:

Check for any outliers in all numeric columns and then analyze carefully, how they should be addressed.

Outliers can be handled by :

1. Replacing them with central tendencies.
2. Removing the outliers in case of extreme out of range.
3. Transforming the data by any of the methods such as dividing the number by larger number or applying log.

```
[30]: # Select numeric columns
numeric_cols = df.select_dtypes(include=[np.number])

# Calculate IQR for each numeric column
Q1 = numeric_cols.quantile(0.25)
Q3 = numeric_cols.quantile(0.75)
IQR = Q3 - Q1

# Define a threshold for outliers
threshold = 1.5

# Identify outliers
outliers = ((numeric_cols < (Q1 - threshold * IQR)) | (numeric_cols > (Q3 +
↪threshold * IQR))).any(axis=1)

# Print outliers
print(df[outliers])

# Remove outliers
df = df[~outliers]
print(df)
```

	carat	cut	color	clarity	depth	table	price	x	y	z	\
2	0.23	Good	E	VS1	56.9	65.0	26160	4.05	4.07	2.31	
8	0.22	Fair	E	VS2	65.1	61.0	26960	3.87	3.78	2.49	
24	0.31	Very Good	J	SI1	58.1	62.0	28240	4.44	4.47	2.59	
35	0.23	Good	F	VS1	58.2	59.0	32160	4.06	4.08	2.37	
42	0.26	Good	D	VS2	65.2	56.0	32240	3.99	4.02	2.61	
...
53882	0.71	Fair	D	VS1	65.4	59.0	219760	5.62	5.58	3.66	
53886	0.70	Good	D	VS2	58.0	62.0	219920	5.78	5.87	3.38	
53890	0.73	Good	E	SI1	57.9	55.0	219920	6.00	5.96	3.46	
53895	0.70	Good	F	VS1	57.8	61.0	220080	5.83	5.79	3.36	
53927	0.79	Good	F	SI1	58.1	59.0	220480	6.06	6.13	3.54	


```

        color_clarity_cut
2          E_VS1_Good
8          E_VS2_Fair
24         J_SI1_Very Good
35         F_VS1_Good
42         D_VS2_Good
...
53882      D_VS1_Fair
53886      D_VS2_Good
53890      E_SI1_Good
53895      F_VS1_Good
53927      F_SI1_Good

```

[6416 rows x 11 columns]

```

        carat      cut color clarity depth table price      x      y      z  \
0         0.23    Ideal      E     SI2   61.5   55.0   26080   3.95   3.98   2.43
1         0.21   Premium      E     SI1   59.8   61.0   26080   3.89   3.84   2.31
3         0.29   Premium      I     VS2   62.4   58.0   26720   4.20   4.23   2.63
4         0.31      Good      J     SI2   63.3   58.0   26800   4.34   4.35   2.75
5         0.24  Very Good      J    VVS2   62.8   57.0   26880   3.94   3.96   2.48
...
53938    0.86   Premium      H     SI2   61.0   58.0  220560   6.15   6.12   3.74
53939    0.75    Ideal      D     SI2   62.2   55.0  220560   5.83   5.87   3.64
53940    0.71   Premium      E     SI1   60.5   55.0  220480   5.79   5.74   3.49
53941    0.71   Premium      F     SI1   59.8   62.0  220480   5.74   5.73   3.43
53942    0.70  Very Good      E     VS2   60.5   59.0  220560   5.71   5.76   3.47

```

```

        color_clarity_cut
0          E_SI2_Ideal
1          E_SI1_Premium
3          I_VS2_Premium
4          J_SI2_Good
5         J_VVS2_Very Good
...
53938      H_SI2_Premium
53939      D_SI2_Ideal
53940      E_SI1_Premium
53941      F_SI1_Premium
53942      E_VS2_Very Good

```

[47527 rows x 11 columns]

Question 7:


Calculate the correlation (Using heat map) between price and all other numeric columns and list them in descending order and identify the highest and lowest correlation?

```
[31]: # Calculate correlation between numeric features
corr = numeric_cols.corr()

# Plot the correlation heatmap
sns.heatmap(corr, annot=True, cmap='coolwarm')
plt.show()

# List the correlation of all features with price in descending order
print(corr['price'].sort_values(ascending=False))

# Identify the highest and lowest correlation
print(f"Highest correlation: {corr['price'].max()}")
print(f"Lowest correlation: {corr['price'].min()}")
```



```
price    1.000000
carat    0.921591
x         0.884457
y         0.865413
z         0.861251
table     0.127118
depth    -0.010666
Name: price, dtype: float64
```

Highest correlation: 1.0

Lowest correlation: -0.01066551076228034

Question 8:

Draw bar plots, visualize and also indicate any insights can be obtained by taking X-axis vs Y-axis as:

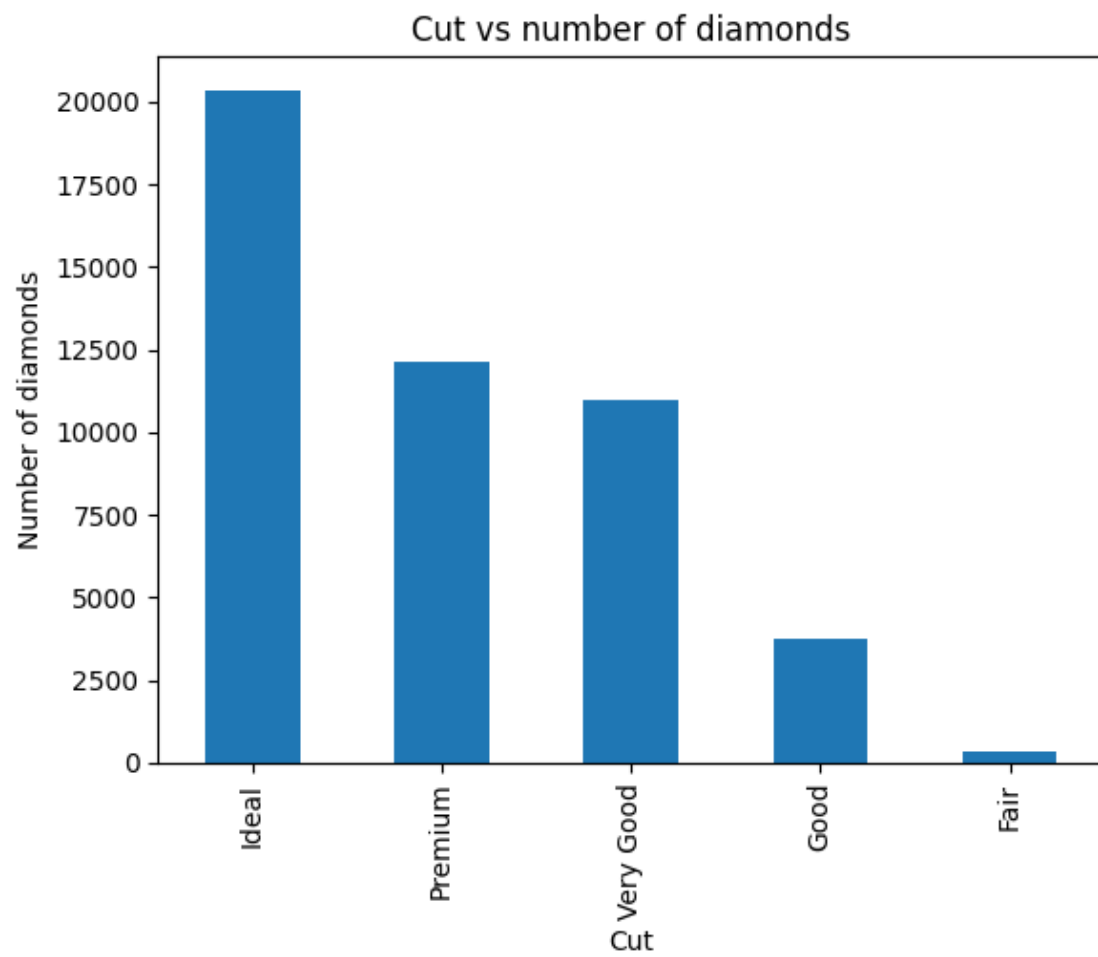
- Cut vs no.of diamonds
- Color vs no.of diamonds
- Clarity vs no.of diamonds


```
[32]: import matplotlib.pyplot as plt

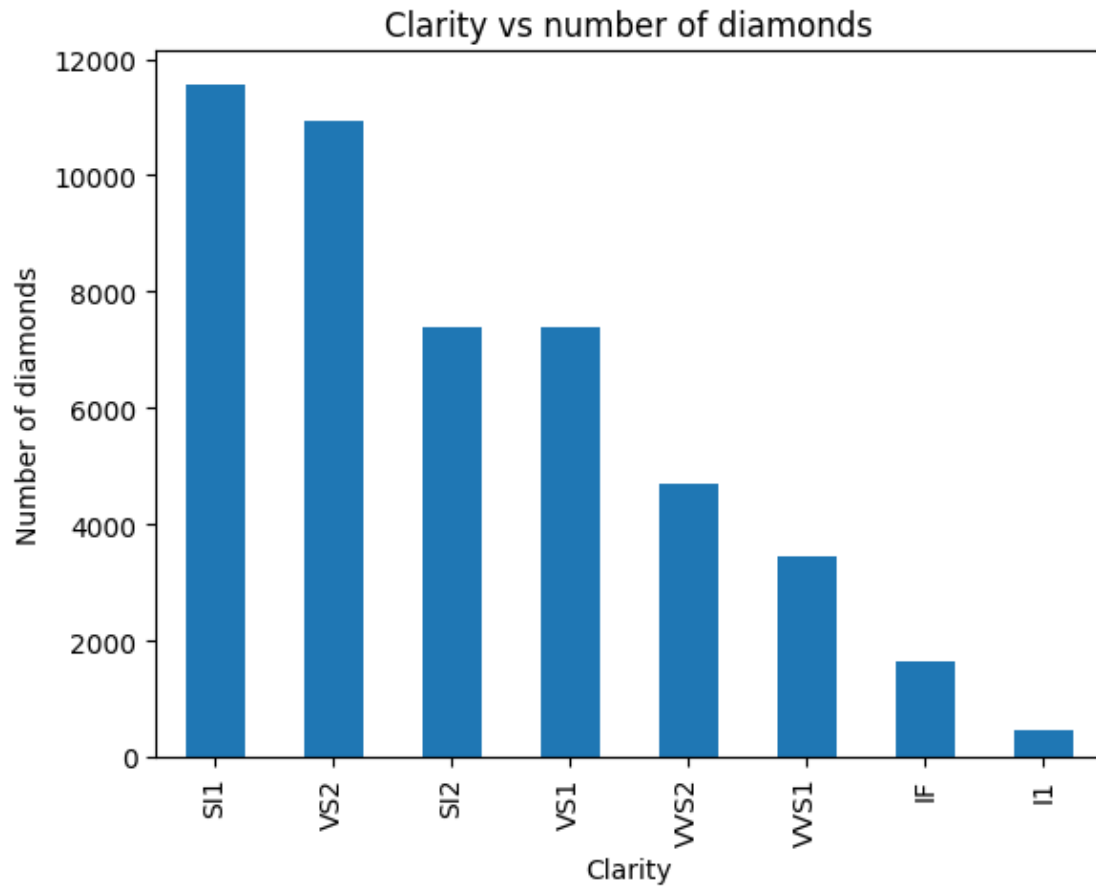
# Cut vs number of diamonds
df['cut'].value_counts().plot(kind='bar')
plt.title('Cut vs number of diamonds')
plt.xlabel('Cut')
plt.ylabel('Number of diamonds')
plt.show()

# Color vs number of diamonds
df['color'].value_counts().plot(kind='bar')
plt.title('Color vs number of diamonds')
plt.xlabel('Color')
plt.ylabel('Number of diamonds')
plt.show()

# Clarity vs number of diamonds
df['clarity'].value_counts().plot(kind='bar')
plt.title('Clarity vs number of diamonds')
plt.xlabel('Clarity')
plt.ylabel('Number of diamonds')
plt.show()
```





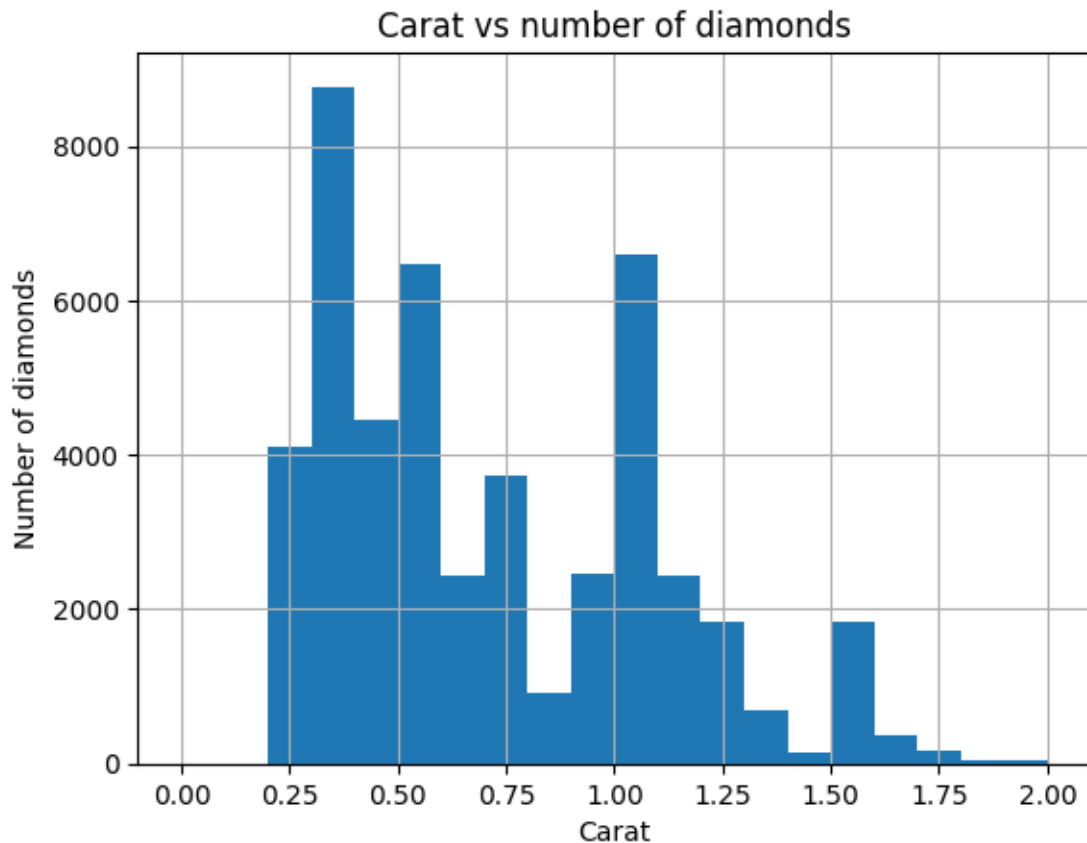


Question 9:

Draw a histogram where X-axis-> carat with interval size 0.1 and
Y-axis-> no. of diamonds? and comment on it

```
[33]: import numpy as np

# Draw a histogram
df['carat'].hist(bins=np.arange(0, df['carat'].max() + 0.1, 0.1))
plt.title('Carat vs number of diamonds')
plt.xlabel('Carat')
plt.ylabel('Number of diamonds')
plt.show()
```



The histogram will show the distribution of diamonds across different carat sizes.

Observations:

- The most common carat sizes among the diamonds.
- Whether the distribution is skewed towards smaller or larger carat sizes.

Question 10:

Draw a normal probability plot on X or Y or z? Based on the shape and trend of the plot? Is any conclusion can be drawn, if yes what it is?

```
[34]: import scipy.stats as stats
import matplotlib.pyplot as plt

# Assuming df is your DataFrame

# Draw a normal probability plot for 'x'
stats.probplot(df['x'], plot=plt)
plt.title('Normal probability plot for x')
plt.show()
```

```
# Draw a normal probability plot for 'y'
stats.probplot(df['y'], plot=plt)
plt.title('Normal probability plot for y')
plt.show()

# Draw a normal probability plot for 'z'
stats.probplot(df['z'], plot=plt)
plt.title('Normal probability plot for z')
plt.show()
```

